

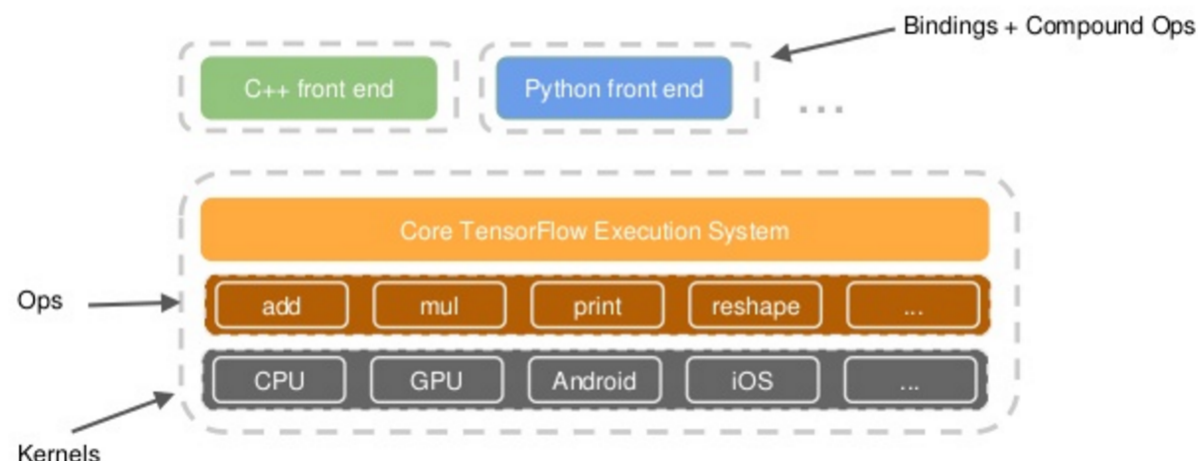
# Deep Learning in TensorFlow for R

Chenglin Zhang  
02/16/2017



# TensorFlow

- Introduced Nov 2015 from Google Brain team
- Deep learning on CPUs and GPUs
- N-d array (tensor) computation on stateful dataflow graph (flow)
- C++ kernel
- Python binding
- Android, iOS  
Raspberry Pi
- Apache 2.0 license
- Open sourced at





<https://www.tensorflow.org/>

<https://github.com/tensorflow/tensorflow>

# Highlights

- Top 15 most popular GitHub repository by stars (passed Linux)
- 1 million+ binary downloads
- Used in teaching university classes  
Toronto, Berkeley, Stanford, ...
- Used in companies  
Google, DeepMind, OpenAI, Twitter, Snapchat, Airbus, Uber, ...
- May 2016, Tensor Processing Unit (TPU)
- Feb 2017, TensorFlow 1.0 release
  - Production ready
  - High level Keras, tflearn interfaces
  - XLA (Accelerated Linear Algebra) experimental compiler
  - Java, Go bindings
  - tfdbg debugger





# TensorFlow for R

TensorFlow R Binding Nov 2016 release from Rstudio

<https://rstudio.github.io/tensorflow/>

<https://github.com/rstudio/tensorflow>

- Tensor
- Variables
- Constants
- Weights
- Biases
- Activation Function
- Loss Function
- Stochastic Decedent Differentiation
- Optimizers
- Graph
- Session
- Input Feeds
- Training
- Tests
- Output Fetches
- Tensorboard Visualization

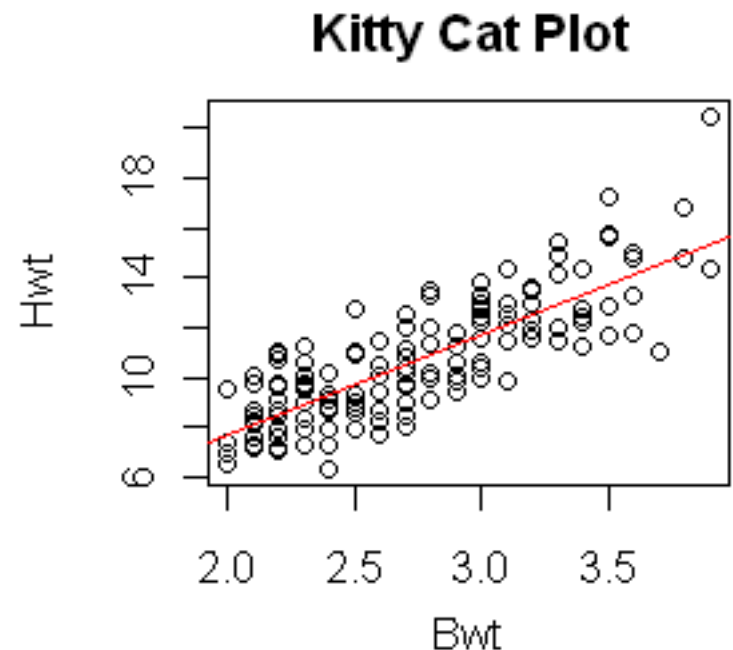
# Linear Regression

- R Jupyter Notebook #1: Demonstrate key concepts of TensorFlow for R, with a sample R cats dataset

```
> library(MASS)
> data(cats)
> mysample <- cats[sample(1:nrow(cats), 10, replace=FALSE),]
> mysample
```


	Sex	Bwt	Hwt
6	F	2.1	7.6
36	F	2.5	10.9
68	M	2.5	8.8
9	F	2.1	8.3
41	F	2.7	10.2
86	M	2.7	12.0
21	F	2.3	8.4
56	M	2.2	9.6
91	M	2.8	11.4
1	F	2.0	7.0

```
> |
```




# Convolutional Neural Networks


- R Jupyter Notebook #2: Demonstrate hand-written digit recognition in convolutional neural network, with MNIST dataset (10 classes)



# VGG 16 Pre-trained Model

- R Jupyter Notebook #3: Demonstrate image recognition in convolution neural network VGG 16, with an imagenet pre-trained model (1000 classes)





# Learning Workflow

- Prepare the data

Data clean up, data normalization, handling missing values

- Select or design a learning model

Convolutional neural networks, recursive neural networks, long short-term memory, deep belief networks, auto-encoders, sequence-to-sequence, deep-q learning, generative adversarial networks, etc

- Choose hyper-parameters

Variable initialization, learning rates, dropouts, epochs, batch sizes, etc. Methods such as random search and grid search can be used to initialize these parameters.

- Train

Fit the model to the data by minimizing losses or cross entropy, by stochastic gradient decent methods or extensions

- Monitor training performance

Maximize training and validation accuracy while avoiding over-fitting to the training data

- Prediction

Deploy and apply trained model to predict



# Accuracy Challenge

2011



**26% errors**

humans



**5% errors**

2016



**3% errors**

# TensorFlow for R Installation

- TensorFlow Installation

[https://www.tensorflow.org/get\\_started/os\\_setup](https://www.tensorflow.org/get_started/os_setup)

# Example: Ubuntu/Linux 64-bit, GPU enabled, Python 3.5

```
$ export
```

```
TF_BINARY_URL=https://storage.googleapis.com/tensorflow/linux/gpu/tensorflow_gpu-0.12.1-cp35-cp35m-linux_x86_64.whl
```

```
$ sudo pip3 install --upgrade $TF_BINARY_URL
```

- TensorFlow for R Installation

<https://github.com/rstudio/tensorflow>

# Example: R with Tensorflow on Python 3

```
Sys.setenv(TENSORFLOW_PYTHON="/usr/bin/python3")
```

```
devtools::install_github("rstudio/tensorflow")
```

- Jupyter R Kernel Installation

<https://github.com/IRkernel/IRkernel>

```
install.packages(c('repr', 'IRdisplay', 'crayon', 'pbdZMQ', 'devtools'))
```

```
devtools::install_github('IRkernel/IRkernel')
```

```
IRkernel::installspec()
```

